

# Knowledge of the U.S. Social Sciences, 1888–1922

*Brooks Ambrose*

*2018-11-06*



# Contents

<b>Knowledge of the U.S. Social Sciences, 1888–1922</b>	<b>11</b>
<b>1 Introduction</b>	<b>13</b>
1.1 Topics $\stackrel{?}{=}$ Ideas . . . . .	15
<b>2 Prior Work</b>	<b>19</b>
2.1 Information . . . . .	19
2.2 Meaning . . . . .	19
2.3 Communication . . . . .	20
<b>3 Disciplines</b>	<b>21</b>
3.1 JSTOR Journals . . . . .	22
3.2 Network Mode Projection . . . . .	22
3.3 Network Community Detection . . . . .	24
3.4 Social Science Journals . . . . .	27
<b>4 Data</b>	<b>31</b>
4.1 Sampling . . . . .	31
4.2 Units of Analysis . . . . .	31
<b>5 Topics</b>	<b>33</b>
5.1 How many topics? . . . . .	34
5.2 Model selection . . . . .	35

<b>6</b>	<b>Web of Knowledge</b>	<b>37</b>
<b>7</b>	<b>Four Fables</b>	<b>39</b>
7.1	Method <i>or</i> The Tortoise and the Hare . . . . .	39
7.2	Ontology <i>or</i> The Fox and the Hedgehog . . . . .	40
7.3	Plan of the Study <i>or</i> A Three-Legged Table . . . . .	41
<b>8</b>	<b>History of Ideas</b>	<b>43</b>
8.1	Antebellum Social Research 1783-1865 . . . . .	43
8.2	Postbellum Social Research 1866-1918 . . . . .	45

# List of Tables

3.1	JSTOR Journal Counts . . . . .	24
3.2	JSTOR Social Sciences Journal Counts . . . . .	27
4.1	Filtering due to Data Management . . . . .	31
5.1	Kurtosis Permutation Test . . . . .	34
7.1	Fable Table . . . . .	41



# List of Figures

1	Explore your options! . . . . .	9
3.1	Mode Conversion on a 300 Edge Random Sample of the JSTOR Journal Label Network . . . . .	23
3.2	Discipline Network (Residual Latent Space controlling for Sociality) . . . . .	26
3.3	Periods in the Growth of the Number of Social Science Journals in the JSTOR Archive . . . . .	27
3.4	Decennial growth in number of PhD degrees conferred in the U.S. . . . .	28
3.5	Number of PhDs conferred in the United States per Social Science Journal . . . . .	29
5.1	Distribution of K by convex hull . . . . .	34
5.2	Significant Counts of K . . . . .	35
6.1	K-clique Community Island Plot . . . . .	37
7.1	Ontological Approaches to Topics . . . . .	42





# Getting Started

Dear reader,

Welcome! This study is available as a website, <https://brooksambrose.github.io/portfolio>, and as a PDF document downloadable from the website. Both are great ways to read the study. The PDF makes for a quicker read, while the website offers additional interactivity in figures and tables that will help you dive more deeply into the exhibits.



Figure 1: Explore your options!

At the top of the web page please notice a toolbar where you can:

- Show and hide the table of contents
- Search the document
- Adjust font and display settings
- View the underlying code at [GitHub.com](https://github.com)
- Download the PDF version

I hope you enjoy the study, and please feel free to report bugs, comment, and collaborate at the issue tracker of the GitHub

repository.

Best,

Brooks

# **Knowledge of the U.S. Social Sciences, 1888–1922**

## **Abstract**

Knowledge development of journals is measured as the change in topic prevalence over time.

## **Keywords**

sociology of knowledge, topic modeling, history of social science



# Chapter 1

## Introduction

What were the ideas that predominated in the social sciences at their formation as professions in the postbellum United States? What was the course of their development over a generation of scholarship? In this study I will answer these questions inductively through a reading of the original journals in each discipline. Though the goal is substantive, the methodological challenges of consuming a large quantity of text will feature importantly in the story that unfolds. Along the way I will demonstrate the usefulness of the computational *distant reading* that is being explored in the humanities and how it can be combined with traditional textual analysis for social science purposes. While controversial in humanistic circles that emphasize the primacy of the reader's novel interpretive work when consuming text, distant reading fits comfortably within a social science epistemology that aims to achieve an objective description of intellectual history. Indeed, computational methods offer a useful backstop to the idiosyncrasy of a particular person's reading of history.

Computational textual analysis promises to automate a particular slice of what hermeneutic methods accomplish. Hermeneutics claims that through historical methods it is possible to reconstruct the interpretive context of texts such that they can be understood in the same way that contemporary historical actors understood them. Establishing such context is a laudable yet arduous feat of historical research to uncover the social and intellectual milieu of a particular text. This is the gold standard approach, but one that restricts the field to specialists with the training and resources necessary for the undertaking.

Computers cannot study history in this way. What they can do, however, is mine source material for limited kinds of contexts. The kind I am concerned with below are the *historical vocabularies* that writers used to construct texts in historical time. Vocabularies are glyphs without grammar; they do not mean anything, but nothing meaningful can be said without them

in the present or in the past. They are the mediated form of language, and in communicating with each other historical actors leave traces that survive perfectly in time so long as texts themselves survive.

While computers cannot read meaning in texts, and can barely recognize it, they are almost as good as humans at recognizing the glyphs of texts, and vocabularies are nothing but glyphs. What computers lack in smarts, they make up in speed and memory. The quantitative scale of their recognition makes for a qualitative shift because vocabularies can be enumerated across immense corpora of texts. Immense, at least, by human standards as there are limits to even computer memory and speed. Yet such enumeration of texts into objective historical categories; this is a profound resource for the intellectual historian. That one could begin a reading with such context would be a transformative research tool. Vocabulary enumeration, by which I mean simply the counting and classifying of texts according to the vocabularies they contain, invites a population studies approach to intellectual history. Where sense-making is driven by comparisons, a reader's arbitrary combination of texts is guaranteed to lead to anachronism. But if we can know that texts are relevant to each other without knowing why, we have done some small amount of hermeneutic work by supplying texts as historically correct context to each other.

And even going so far as abandoning the project of reading texts in a historically correct way, vocabulary enumeration can still lend objectivity to a novel construction, a productive anachronism, of textual meaning. Because vocabularies, the problems solved by computers, are mathematically, algorithmically, or stochastically determined, they may provide an immutable description of corpora that, like a map, enables individual and collective exploration within a common framework. Such maps may become the parameters of interpretive methods, which we may use to surface and control some of our subjectivity.

This at least is the rationale for what follows. I begin with a discussion of intellectual history of two social sciences, anthropology and sociology, in the United States. I take a coarse view of national history as the history of wars because of their downstream effects on government activity and institutional investments. The first period is between the end of the American Revolution (1783) and the end of the American Civil War (1865) and is the national context for the origin of U.S. anthropology. The second period is after the Civil War until the end of World War I (1918) and is the context for the origin of U.S. sociology and of modern U.S. higher education generally. Wars of territorial expansion are waged regularly during both periods against native peoples and rival colonial empires, and social research was always recruited to solve attendant problems of population and to provide rationales for the relationships with and understandings of conquered or would-be conquered people.

I use intellectual histories of anthropology to characterize the antebellum period, and the same for the postbellum period including sociology. The most important journals in each field date from the postbellum period, and the appearance of each

is implicated in the project of professionalization for each discipline. The 1920s marked the end of war with the last of the militating American Indian tribes, and a reckoning with the darkest sides of industrialization laid bare by WWI. Social research had by this time completed a shift from colonial to industrial problems and enjoyed a golden decade of development as a profession, punctuated by the next great historical crisis in the Great Depression. With the 1920s begins the adolescence of social research, which is beyond the present scope. This study is of its childhood, which ends with the Great War. I however draw the study out until 1922 because it is the end of the public domain in U.S. copyright, to aid in the reproducibility of the analysis and so that all readers may recover the texts in question without difficulty.

## 1.1 Topics $\stackrel{?}{=}$ Ideas

The strategy of the study occurs in four steps.

1. Sort text into categories of similar vocabulary.
2. Describe the vocabularies that define category membership.
3. Describe vocabulary prevalence across time and discipline.
4. Validate category contents by a traditional qualitative reading of texts.

I will spend considerable effort on solving the problem presented by step 1, as here everything depends on the computational methods employed. Steps 2 and 3 are straightforward given a successful mathematical model of texts. Step 4 is seldom attempted, and may be the hardest of all, because it is here that machine and human learning must be integrated. If I am successful, if through these steps I may operationalize the notion of cultural meaning or cultural logic as conformity to vocabularies, then I believe a new horizon of intellectual scholarship is possible. If on the other hand I find that machine-learned vocabularies do not correspond to human-learned understandings of the texts drawing on those vocabularies, then the discovery will be negative, that distant reading is not a scientific, historical, or hermeneutic method, but rather a toy at worst and a best new humanistic method of reading texts *de novo*.

The mathematical tool I will rely on in step 1 is called topic modeling, which refers to a variety of computational approaches to text data that blur the distinction between qualitative and quantitative analysis. The topic model paints a lexicographic picture of texts, analogous to the demographic picture gained by a civil census survey of cities and towns. To a topic model, texts are merely collections of terms (usually words) that are counted to create the so-called “bag of words” description of a text. In

the same way that a census reduces communities to counts of the names of people who live in them, topic modeling reduces texts to the frequency of word choices in texts, to their diction or vocabulary. Just as a census of people fails to capture the nuanced interactivity of human settlements found in their culture, politics, and economic activity, the topic model washes away the meanings and intentions behind the words that are enumerated.

A population census would not be very helpful were it only a count of the names of respondents, and of course the really helpful data derive from the demographic and economic survey attached to the name. Text data do not usually come with such a collection of rich covariates, yet nevertheless topic models promise to discern helpful patterns from counts alone. The trick behind the estimation of a topic model is that it attempts to learn the demographic information (topics) without asking, by merely looking at how the names alone (terms) are distributed across geographies of interest (texts). If it can keep its promise, a topic model applied to census data might recover the cultural patterns latent in the distribution of names. It might, for instance, learn different groupings of names that in turn correspond to markers like age, race, national origin, or gender, so long as membership in those categories was related to geography. It might, for instance, successfully separate a category of Hmong names out from among the names of all people living in St. Paul because the non-Hmong names appeared in other regions where no Hmong names appeared.

To call the category of names “Hmong” requires an interpretation of the model, which by itself is just lists of names. This is the work of step 2, and requires a little bit of shoe leather by trying to make sense of what a list of names refers to. Here reading texts is like a census taker knocking on a door, and a topic model’s latent analysis saves on this effort. Sometimes bringing domain knowledge to bear on the list itself will suggest a category label, but often choosing a small sample of texts as exemplars of the category. Still this requires much less shoe leather than a traditional qualitative analysis in which each text is studied directly. Of course the census is much more informative because it asks about demographic categories directly thereby avoiding the need for a latent analysis. In domains where rich covariates are not yet available or are prohibitively expensive to acquire, latent analysis provides promising clues of patterns that already exist. What is even more interesting, and something that might surprise even census analysts, is when latent categories do not correspond to known survey items. In either event the power of topic modeling for inductive analysis is to reveal structure in how names hang together that was hidden.

Even without conducting the second labeling step, in step 3 it will already be possible from the output of the model to inspect the distribution of topics across available covariates, especially time. These are the patterns that will help validate the topic models against what is already known about intellectual history. For instance, the power of institutional and generational



change may well be apparent in the historical distribution of topics. This step leads naturally into step 4 by suggesting anomalies that can only be explained by a closer look at the texts, the chore that the entire preceding analysis punts on. In step 4 we learn either that our understanding of history was wrong, or that our topic model was wrong, and there may be no method other than one's judgement to decide.

In the next section, before we delve into the statistical and computational nuances of topic models, I will spend some time developing a few themes to help organize the blending of quantitative and qualitative methods invited by topic modeling in particular and computational text analysis generally.



## Chapter 2

# Prior Work

### 2.1 Information

Understanding differences in the ontological status of the “topic” concept is a good way to begin to understand how this method of analysis is used by researchers.

Analysts have conceptualized the use of topic models in very different ways. Some researchers treat topics as useful for a particular purpose and not as true descriptions of real phenomena. Topics as information enhances the ability to search for relevant documents or statistical trends in otherwise unwieldy corpora as a time-saving alternative to manually reading large collections. (Boyd-Graber et al., 2017) Empirical problems, used as demonstrations of statistical techniques, have included

This is the “needle and haystack” approach favored by computer and information scientists who tend not to be interested in theoretical interpretations beyond the statistical definitions of topics.

### 2.2 Meaning

Other researchers instead grant topics ontological status, and these can be divided into three types. Most ambitiously, topics may be treated as representing categories of thought. Latent semantic structure latent semantic structure (?)

## 2.3 Communication

representational style (Grimmer, 2016) frame (DiMaggio et al., 2013)

## Chapter 3

# Disciplines

Computational text analysis requires that text corpora be transformed from a human to a machine readable format. Several efforts to digitize paper archives have made historical research designs possible, notably the Google Books project, HathiTrust, and ITHAKA JSTOR archive. Digital storage devices like the portable document format (PDF) have also enabled texts to be represented in both a digital version and as a reasonable facsimile of paper originals. Reasonable, we should say, for most sociological purposes, but not for other historical questions where materiality of culture is important. (Schreibman, 2014, 149)

Digital archives make research into the production of culture difficult, precisely because they misrepresent several aspects of the means of production. Because researchers should be mindful that digitization of texts abstracts some qualities of texts and renders many others invisible. The importance of physical space and material qualities of libraries is illegible when working with digital archives, while the verbal content of texts is highlighted. We must keep in mind that we are not viewing what historical actors saw. Digital texts are almost perfectly fungible, while, variability in historical texts. We are liable, for instance, to underestimate the search costs to locate texts, and the fungibility of texts themselves.

There are reasons, however, to believe that digital text archives provide not just a useful but an historically valid abstraction from the material texts. If we want to understand how an individual scholar understood a particular text, better to have her personal copy, margin notes and all. Yet how would that scholar have treated the text as a cultural item? She would abstract her own copy to a format credibly held in common, the more anisepotically clean version that we see in digital archives. These are the ghosts of the texts, so to speak, but they are what would be left when all idiosyncracies were removed, the version that one would assume colleagues thought of when declaring that text publically.

This is by way of saying that the texts I compile below are not the same that were read by the historical actors under consideration. They are the texts that historical actors would assume their contemporaries were reading, that is, the sanitized, fungible, original published form of the text. By getting at these texts, we are getting at the real historical infrastructure for scholarly communication.

The optical character recognition that computers require in order to store text digitally depends critically on the hard work of creating quality scans of journal archives. JSTOR has done a comendable job of this. Next we will describe what the JSTOR archive has to offer.

### 3.1 JSTOR Journals

We rely on the JSTOR digital archive which gives access to optical scans of historical journals. JSTOR provides a title list of their journal coverage (JSTOR, 2018). The coverage of journals in the archive is very complete for those journals chosen for the database. As of this writing JSTOR contained 4,162 different journal titles and 2,706 journals from 1,121 different publishers. The different journal counts are due to some journals changing titles at least once.<sup>1</sup> These journals are organized into nine superdisciplines under which may be found 70 disciplines.

Most journals are given more than one discipline label, and the superdisciplines are not marked as such in the database creating some redundancy. For instance, a journal labeled as “Sociology” will also be labeled as “Social Sciences”. Most academics will be familiar with whether a label is for a superdiscipline or a subdiscipline, yet for outsiders or for skeptical insiders, the only clue is in the frequency with which a label is applied. Counting labels, however, does not unambiguously place a journal in one discipline or another because journals may bear multiple labels, even multiple superdiscipline labels.

To assess the size of the disciplines and to disentangle their hierarchies it will be helpful to have a mutually exclusive labeling scheme that draws on the JSTOR curators’ judgement while simplifying it.

### 3.2 Network Mode Projection

I rely on network methods to accomplish this labeling in a data driven and reproducible way. First I projected the journals as a bimodal network with journals in the first mode bearing ties to discipline labels in the second mode. I then projected

---

<sup>1</sup>To avoid overcounting, title histories are collapsed into their most recent record, meaning all subsequent counts are out of 2,706. Even though we might expect disciplinary identity to change over time, JSTOR discipline labels do not vary within title histories. Two journals—Haldane Bulletin and Scientific American Mind—lacked any discipline labels and are excluded from tabulations.

the bimodal network into two unimodal networks, one where journals are connected by ties equal to the number of discipline labels they have in common, and another where labels are tied by the number of articles carrying both labels. Call each of these unimodal networks, the (*journal-label-journal*) journal network and (*label-journal-label*) label network, a facet of the original bimodal network.

Figure 3.1 illustrates the effects of network mode projection on a random sample of 300 edges. The first panel illustrates the bimodal network where journals are yellow dots and labels blue dots. As an artifact of sampling, most journals here are shown tied to only one label. In fact this is never the case in the full network; as each journal has at least one discipline and one superdiscipline label the minimum number of labels is two, the median case accounting for 53.6 percent of journals. The most labels any journal bears is 10, but these are outliers with most journals bearing only a few labels.




Figure 3.1: Mode Conversion on a 300 Edge Random Sample of the JSTOR Journal Label Network

It is worth noting a few features of unimodal projections. First, unimodal projections will always be made of up of overlapping cliques. Take the journal facet; each journal bearing a particular label will be tied to every other journal with the same label, and together they will form a maximal clique, a subnetwork where all possible ties exist. Such cliques grow nearly exponentially, as each additional journal with the same label joins the clique and adds a number of ties equal to the former size of the clique. Practically, very common labels like “Social Sciences” can easily dominate a network. Here the weighting of edges becomes important; if two labels overlap because some nodes bear both labels, then within the intersection of the two groups the ties may “weigh more” by adding the contribution of each label separately. Using methods that take edge weights into account is a good way to ameliorate the exponential influence of popular labels.

Second, though unimodal facets of a bimodal network represent the same data, each may have different characteristics

especially in the common case of a large population imbalance between modes. Here we have 34 times as many journals as labels and each journal sends multiple ties. This degree imbalance between the two modes may mean that one facet is more dense than its inverse. Density is the proportion of actual ties out of all possible ties. An imbalance may be observed in the sample network where the journal facet has many dense and interconnected cliques and where the label network appears to be mostly made of isolated labels save for the few larger components. Here the journal facet is seven times more dense than the label facet.

While each of the isolates in the label facet represents a larger or smaller number of journals (se the free floating cliques observable in the journal facet) they supply no information about interdisciplinarity. Practically, analyzing the journal network highlights the population sizes, while analyzing the label network highlights interdisciplinarity.

In the case of our full network, the potential imbalance in degree distribution between facets happens to be offset by the population imbalance itself. The densities in the full journal and label facets are comparable, 26 and 27 percent respectively, meaning that analysis will not merely hinge on which facet is analyzed.

### 3.3 Network Community Detection

With our facets in hand we may now tackle the problem of partitioning the JSTOR journals into mutually exclusive disciplines. I define disciplines as clustering within each of the two facet networks. I conducted community detection—a set of network methods that classify nodes according to the dense subnetworks they occupy—to cluster each network and reveal patterns in JSTOR’s labeling policy. (See appendix ? for technical details.) The goal was to confirm the expectation of which labels counted as superdisciplines, or to learn a different classification scheme if the results were surprising.

Table 3.1: JSTOR Journal Counts

Superdiscipline	Learned	Pct	Given	GPct
Social Sciences	752	27.8	916	33.9
Humanities	666	24.6	688	25.4
Area Studies	355	13.1	499	18.4
Science & Mathematics	311	11.5	360	13.3
Business & Economics	264	9.8	284	10.5
Arts	244	9	292	10.8
Law	84	3.1	132	4.9
Medicine & Allied Health	30	1.1	52	1.9
Total	2706	100	3223	119.1



The results in Table 3.1 are achieved by partitioning the *journal* network into the detected clusters, and then by assigning each cluster the name of its most frequent label. The learned labels refine the superdisciplines by making them mutually exclusive, as is apparent by the sum of the given labels exceeding 100%. The cluster labels comport with a commonsensical understanding of superdisciplines. Social Sciences journals predominate due to JSTOR's initial focus in that area, even without counting economics among them. Business & Economics marks the transition from the larger academic journal space to the smaller professional space.

Clustering of the label facet presents a similar view. Recall that each facet represents the same data, the difference being whether a journal or a label is represented as a node or an edge. The larger population of journals is easier to partition it into a greater number of subpopulations. Converseley, because there are far fewer labels than journals, we would expect the clustering to be less granular for the label network than for the journal network. In fact there is only one fewer clusters, with the difference being Law.

Figure 3.2 shows a graph of the relationships among labels, where again the strength of ties is equal to the number of journals bearing both labels. Again community detection partitions the network into its most dense subregions. Here the label with highest degree is assigned to the category and then omitted as a node. Astronomy, an isolate that never accompanies another label, is also omitted from the graph, resulting in four superdiscipline categories.




Figure 3.2: Discipline Network (Residual Latent Space controlling for Sociality)

Here some of the granular categories are collapsed. The humanities includes arts, as we might expect, but also area studies, which one might have classed with the social sciences, but which bear stronger ties to cultural studies like music, folklore, religion, and language and literature. Law and medicine and allied health are grouped with the social sciences, and business and economics is maintained as separate field due merely to the attachment of three professional fields—development studies, management and organizational behavior, and marketing and advertising—to their parent disciplines business and economics (not to be confused with the separate and omitted label “business and economics”), which are themselves strongly tied to the social sciences.

Though graph layouts are imperfect and should not be overinterpreted, the global features of facing within clusters do indicate the disciplines that straddle boundaries. On the border between the social sciences and science and mathematics are the social sciences dealing most with the physical problems of space, health, and technology. On the edge of the humanities and social sciences are history, philosophy, and anthropology.

### 3.4 Social Science Journals

The journals within social science cover five different subdisciplines.

Table 3.2: JSTOR Social Sciences Journal Counts

Subdiscipline	N	Pct	Labeled	LPct
Archaeology	256	27.9	115	12.6
Political Science	215	23.5	180	19.7
Education	196	21.4	174	19
Sociology	160	17.5	145	15.8
Anthropology	46	5	89	9.7
Population Studies	22	2.4	27	2.9
Geography	18	2	32	3.5
Transportation Studies	3	0.3	7	0.8
Total	916	100	769	84




Figure 3.3: Periods in the Growth of the Number of Social Science Journals in the JSTOR Archive




Figure 3.4: Decennial growth in number of PhD degrees conferred in the U.S.




Figure 3.5: Number of PhDs conferred in the United States per Social Science Journal

This period represents one of stable growth, as the size of the field grows with the number of players on it. Between 1888 and 1922 there tended to be about seven new PhDs in the U.S. for every social science journal even as each population grew year over year. These growth patterns begin to diverge around 1920 as a decades long acceleration of personnel begins, relatively slowly between 1920 and 1960 at an average acceleration rate of 22 PhDs per journal per year, and then quite precipitously in the 1960s at an average acceleration rate of 121.



# Chapter 4

## Data

Every record for every journal was downloaded manually, including front and back matter, articles, and book reviews.

### 4.1 Sampling

Table 4.1: Filtering due to Data Management

step	doc	pag	par	sen	tok	ter	lem
imported	100	100	100				
cleaned	99.27	98.21	67.51				
tokenized	99.27	98.21	67.51	100	100	100	
preprocessed	99.27	98.01	67.35	91.38	42.21	35.74	100
sampled	1.84	1.56	1.17	1.43	0.62	4.95	20.86
100							

### 4.2 Units of Analysis

Conventionally researchers feed entire documents into the construction of term frequencies. This method treats any term in a document as being related to any other term by the same degree. The goal of any topic mixture model algorithm is to sift these terms into different topic categories basically by looking for clues across documents; a topic can be “seen” in a particular document to the extent that other documents include that topic and *other* topics different from the focal article, so that the intersection of terms reveals the topic. But a much simpler assumption to reduce the attendant noise within a document is

to merely feed lower level syntactic structures—paragraphs and sentences—to the algorithm. We will see that doing so greatly improves the usefulness of discovered topics.

The irony of this approach is that while topics become more clear as documents become shorter, the assignment of any particular shorter document to a topic is murkier due to the smaller word count.

Long documents will contribute more text to the corpus, but this is fair as they make up more of the population of text. Thus a simple random sample will allow better descriptive statistics. I sampled at the paragraph level because.



## **Chapter 5**

### **Topics**

The modeling objective is twofold, to sort text into categories of similarity, and to describe the qualitative content that defines the category membership. In this way we may operationalize the notion of cultural meaning or cultural logic as the rules of category classification. reduce expressions as instances of a latent category of expression.

## 5.1 How many topics?




Figure 5.1: Distribution of K by convex hull

Table 5.1: Kurtosis Permutation Test

level	e	se	l99	u99	$P(e \leq 0)$
doc	-0.0932	0.1149	-0.3682	0.2252	0.7948
par	-0.1125	0.1206	-0.3999	0.2185	0.8257
sen	0.0118	0.2304	-0.5078	0.6471	0.4973




Figure 5.2: Significant Counts of K

## 5.2 Model selection



## Chapter 6

# Web of Knowledge




Figure 6.1: K-clique Community Island Plot



## Chapter 7

# Four Fables

What were the dominant ideas at the genesis of U.S. social research? How would we know them if we found them? I consider two sets of approaches, the first methodological (how?), and the second ontological (what, why?), which will help orient us to the plan of the study. I will briefly describe a version of the familiar distinctions among qualitative and quantitative methods on one hand and nomathetic and idiographic ontologies on the other. What is new is the question of whether computational text analysis may blur the lines between these classic social science epistemologies. In answering this question I will organize some existing literature into four categories of research.

### 7.1 Method *or* The Tortoise and the Hare

Methods are procedures for arriving at results, and they act by exposing assumptions to empirical observations. They provide a source of influence on arguments that breaks the circularity of their reasoning. Methods differ in the amount of material exposed to the observer and the speed of its exposure. The difference between human and machine learning concerns both, but speed is especially salient.

“You may deride my awkward pace,

But *slow and steady* wins the race.”

– Robert Lloyd (1762, 38)

Lloyd’s version of Aesop’s contestants correctly describe the consequences of haste.

First consider a humanist, a historian, a scholar with sense enough to read primary source material. How would she proceed to conduct an intellectual history? Slow and steady, a tortoise would identify diverse documentary sources allowing her to collect the names of important people and organizations and learn what she could of their biographical facts and event timelines. She would then read the scholarship both produced and consumed by these important actors. The identification of ideas would be the most difficult of her tasks, not just because reading takes time and effort (and there would be much of it), but because ideas exist in the minds of people and leave no direct empirical trace. The interpretation of writing in a historical context would suggest candidate ideas; their prevalence across time and place would indicate their historical importance. It would be tedious work requiring intelligence and patience.

Next consider a sociologist, a computer scientist, a librarian awash in books with no time to read them. She would identify a convenient source and ask a computer to do the rest. Hasty, impudent, and lacking the tortoise's fortitude and patience, a hare uses a mental prosthetic to achieve and perhaps exceed the scale at which a humanist can consume documentation. She learns much less than the tortoise because, whereas people can find what they were not looking for, machines can learn only what they are told to learn. She fails in the test of knowledge (of course we know who wins in the end), but at least she fails quickly.

For shorthand, we can refer to these two approaches as Aesop's tortoise and hare, the humanist mode and the computational mode.<sup>1</sup>

## 7.2 Ontology *or* The Fox and the Hedgehog

So we might learn more slowly or less quickly, but what will we be trying to learn, and why? These concerns the question of ontology. Ontologies are categories of being, or more simply, they are the assumptions that answer the (usually implicit) question, "What is this?". Ontologies may be descriptive categories of classification or explanatory causal mechanisms. It is neither possible to describe nor explain without making ontological assumptions.

"The fox knows many things,

but the hedgehog knows one big thing."

– Isaiah Berlin(1953, 1)

Here two of Berlin's creatures will help show two different ontological approaches.

---

<sup>1</sup>The tortoise may just as well be a hammer, and the hare a steam drill, but for the fact that the hammer ended up losing his contest.



Consider a surveyor, an ethnographer, a data scientist. A fox believes that the world is nothing but the facts about it. She sets out to learn something about everything, and in so doing she tends to locate where the action happens to be without having known it was there to begin with. Though she can only skim for surface features, a fox's shallower understanding of many things is usually very helpful. A fox learns where the important, useful, or interesting things in the world are hidden.

Next consider a theorist, a statistician, a case worker. A hedgehog learns something, but not necessarily everything, about something. They know less than a fox, but if they are a good hedgehog then what they know is good enough. A hedgehog experiences the world in filtered fashion, bothering to remember only what contributes to her system or her obsession.

So let us run Aesop's race again, but make it a three-legged race competing in teams. And to add some purpose above crossing the finish line, the party set an objective. The teams must gather fruit to bring back to the table for supper. Who would win? Hare and hedgehog of course, for they knew a big melon would provide more than their share, which they found at a farm stand down the road. The real contest was for second place. Hare and fox set out at once to forage, darting furiously here and there to gather a great assortment of wild nuts and berries. There were many mouths to feed and each morsel was small, so they had to be dedicated to growing a larder. Tortoise and hedgehog heard tale of a durian, the spiked and malodorous king of fruits, and thought it fit well with hedgehog's general motif; but they did not sell them at the country market so they walked into town looking for an importer. They returned at sunset, durian in tow, to find hare and fox laying exhausted atop an impressive heap of produce that hare and hedgehog had clearly been snacking on while they idled the afternoon away. Sadly tortoise and hare never even placed; they had thought it prudent to plant a garden of their own and left for the seed store with no intention of returning to the party. In a few years time they intended to be the life of the potlach.

### 7.3 Plan of the Study *or* A Three-Legged Table

So in four corners of a table, three were full of different produce and in one lay only a promise.

Table 7.1: Fable Table

	fox	hedgehog
tortoise	Historian, gold standard	Historian, gold standard
hare	Historian, gold standard	

"One's own consciousness dances about upon the words like a will-o'-the-wisp."

—Niklas Luhmann (2002, 166)




Figure 7.1: Ontological Approaches to Topics

Text as thought, or as communication. Thought holds that the ideas can be reliably interpreted, perhaps hermeneutically, to recover the mental events or intentions of authors and readers. Ambitious. Easier is to treat texts as communications, as messages, and worry little about their meanings or interpretations. A study of communication is a good foundation for the study of thought, but it is a separate task and the one we undertake here.

## Chapter 8

# History of Ideas

Though the intellectual history of the social sciences begins well before the Civil War, the current epoch of its institutional history, the epoch of professions, becomes possible only in the postbellum period. Before the Civil War, social research was a skilled occupation, and individual researchers found patrons through government and civil institutions supporting pursuits of knowledge, such as the American Philosophical Society (APS) founded in 1743 in Philadelphia. One of the great consequences of war was renewed federal investment in nation building in its aftermath. Some Union congressmen saw the rebellion as a failure of education, prompting the creation of the U.S. Department of Education in 1867 to strengthen nation building. In the postbellum period both anthropology and sociology develop as professions due primarily to the growth of universities as a new context of their activity. Without universities, social research would have remained an occupation in need of clientele. With them, social research develops the resources enabling relative autonomy, self reproduction, and occupational closure, the hallmarks of a profession. This study seeks to examine the period of transition after the Civil War from social research as occupation to social research as profession.

### 8.1 Antebellum Social Research 1783-1865

U.S. nation building had continued since the end of the American Revolution and had enrolled researchers in the projects of westward expansion against native peoples, the consolidation of slave economies against Africans, and the legitimation of the American experiment against European detractors. These were pressing problems to the intellectuals among government leaders at different levels, and they worked to make investments in new knowledge to resolve them. Such new knowledge was

initially an extension of older “theories of man” in theology and enlightenment natural philosophy, which had a foothold in the private education of the American so-called natural aristocracy as well as in urban colonial institutions like the APS serving as meeting places for intellectual elites and scholars. After the British burned the Library of Congress in 1814 Thomas Jefferson famously sold his personal library to Congress to restore it, an illustration that secular arts and sciences were produced and maintained by and under the patronage of private elites.

The “theory of man” in the colonial period was that historical progress played out along a scale from savagery to barbarism to civilization, and that movement in the direction of progress was a function of time and innate capacity. European societies of the Old World had the best endowment of each, for they were very old (seen as continuous with pre-Christian antiquity) and were lead by aristocracies and monarchies representing the highest human capacities. Conversely colonized peoples, and colonial societies themselves, were understood to suffer low endowments of both, developing slowly due to their lack of morality, reason, and aesthetics and having had less time to achieve what little progress they could.

To gain parity with their colonial masters, the American elite altered the theory to emphasize capacity over time. Because all men were equal, the Americas could achieve greatness meeting or surpassing Europe given independence and opportunity. Their masters scoffed at the idea that a colony could ever “catch up” to the development of Europe, with aristocratic families charting their roots into antiquity. American elites, especially Southerners like Jefferson, saw themselves as naturally superior to their countrymen, taking an aristocratic role in their own country, but not requiring the ancient lineages qualifying aristocracy in Europe. The values of the Declaration were egalitarian only in this limited sense of parity across civilizations.

Egalitarianism did not mean parity within civilizations. Colonial theologians, philosophers, and learned elites had long provided the intellectual rationalizations of the domination of women and the environment, and such reasoning was readily exported to the yet unresolved problems of native tribes and African slaves. Paradoxically egalitarian values reinforced rather than ameliorated supremacist ideologies at home. If all men were equal, and natives (or Africans, or women) were clearly inferior, then natives must not be men. The racism of the Old World order in which hierarchy and domination were natural and unproblematic, was also a consequence of egalitarian values but for different reasons. Such racism was indeed more pernicious in America as the concept of equality invited even more ardent “proof” of the inferiority of dominated peoples.

American intellectuals knew the same logic of inferiority could be applied to the American elite who if they could not match the rate of progress of European society would be judged a newly diverging and potentially inferior race of men. Though it did little to ameliorate white male supremacist thinking, the contradictions between the universalistic and egalitarian values

expressed in the U.S. Constitution did occasionally create grist for academic debate leading to some investments in research critical of extant racist typologies. In Europe scholars of comparative languages charted the supposed common origins of European cultures and the time it had taken them to diverge. This inspired Jefferson to patronize ethnological research on the languages of native tribes in America. He had by 1785 amassed enough data to appreciate their diversity, which he took as evidence that New World cultures may be older, and therefore more developed, than commonly assumed by colonizers.

The egalitarian-for-its-time thinking of Jefferson did not contravene racism, but it did provide a different model of racial hierarchy, which could be consequential for government projects. Jefferson thought that if all men are equal, that they share the same human capacity for progress, then native people had merely not had the time to develop and discover progress for themselves. They could and should be taught, especially to abandon hunting in favor of agriculture. For a brief time in the 1840s the Virginia legislature used tax incentives and educational programs to promote intermarriage between male settlers and native women with the goal of accelerating the natural development of native culture. (Patterson, 2001, 9) Such experiments were short lived, as the predictable backlash was that such mixing would risk lowering white culture rather than elevating Indian, and genocidal and segregationist policies won out. Thus while all racial models concluded with the domination of natives by settlers, different models of racial hierarchy sewed policy disputes within colonial leadership.

The role of social research then was to be in service of either government or private political associations. Careers were made for anthropologists in patronage relationships to generate knowledge to aid nation builders.

## **8.2 Postbellum Social Research 1866-1918**

While the roots of anthropology are as old as the republic, sociology did not develop in earnest until after the Civil War.



# Bibliography

Berlin, I. (1953). *The Hedgehog and the Fox: An Essay on Tolstoy's View of History*. Weidenfeld & Nicolson, London. OCLC: 10872907 Citation Key: Berlin1953Hedgehog.

Boyd-Graber, J., Hu, Y., and Mimno, D. (2017). Applications of Topic Models. *Foundations and Trends® in Information Retrieval*, 11(2-3):143–296. Citation Key: Boyd-Graber2017Applications.

DiMaggio, P., Nag, M., and Blei, D. (2013). Exploiting affinities between topic modeling and the sociological perspective on culture: Application to newspaper coverage of U.S. government arts funding. *Poetics*, 41(6):570–606. Citation Key: DiMaggio2013Exploiting.

Grimmer, J. (2016). Measuring Representational Style in the House: The Tea Party, Obama, and Legislators' Changing Expressed Priorities. In Alvarez, R. M., editor, *Computational Social Science: Discovery and Prediction*, Analytical Methods for Social Research, pages 225–245. Cambridge University Press, New York, NY, reprint edition edition. Citation Key: Grimmer2016Measuring.

JSTOR (2018). Title Lists. Citation Key: JSTOR2018Title.

Lloyd, R. (1762). *Poems*. author. Google-Books-ID: t71IAAAAMAAJ Citation Key: Lloyd1762Poems.

Luhmann, N. and Rasch, W. (2002). *Theories of Distinction: Redescribing the Descriptions of Modernity*. Stanford University Press. Google-Books-ID: L0G2bwV9VMCM Citation Key: Luhmann2002Theories.

Patterson, T. C. (2001). *A Social History of Anthropology in the United States*. Berg, Oxford; New York. OCLC: 48551832 Citation Key: Patterson2001Social.

Schreibman, S. (2014). Non-Consumptive Reading. In Segal, N. and Koleva, D., editors, *From Literature to Cultural Literacy*, pages 148–165. Palgrave Macmillan UK, London. Citation Key: Schreibman2014NonConsumptive.